

# Assignment 0

## Code Walkthrough Questions

A page for students to agree on answers to the questions posed in the code walkthrough. Perhaps those of us who are bored or scared will fill it out.

*I've added the answers we came up with. Maybe some others can correct any mistakes we made - Nathan Wilson and Ivor Popovic*

**Question 1: What is the vm system called that is configured for assignment 0?**

**Answer 1:** dumbvm

*Refer to: kern/arch/mips/conf.arch, implemented in kern/arch/mips/vm/dumbvm.c*

**Question 2. Which register number is used for the stack pointer (sp) in OS/161?**

**Answer 2:** \$29

*Refer to: kern/arch/mips/include/kern/regdefs.h*

*Particular line: #define sp \$29 /\* stack pointer \*/*

**Question 3. What bus/busses does OS/161 support?**

**Answer 3:** LAMEbus

*Refer to: kern/arch/sys161/include/bus.h*

**Question 4. Why do we use typedefs like uint32\_t instead of simply saying "int"?**

**Answer 4:** To specify that a value is an unsigned 32 bit int. Different architectures handle 'int' as different sizes. Specifying like this allows us to make sure that it will be a consistent number of registers/memory each time - always 32bits, unsigned, and thus allows for simplifying compatibility across platforms. It also allows us to get more mileage out of a variable that we know will never go negative.

**Question 5. What function is called when user-level code generates a fatal fault?**

**Answer 5:** static void kill\_curthread(vaddr\_t epc, unsigned code, vaddr\_t vaddr)

*Refer to: kern/arch/mips/locore/trap.c*

**Question 6. How frequently are hardclock interrupts generated?**

**Answer 6:** 100 times per second.

*Refer to: kern/include/clock.h*

*Particular lines:*

切换行号显示

```
1 /* hardclocks per second */
2 #define HZ 100
3
```

**Question 7. How many characters are allowed in an SFS volume name?**

**Answer 7:** The maximum is 32 characters including the null termination character, so only 31 characters are allowed in the name.

Refer to: `kern/include/kern/sfs.h`,

Particular line:

```
#define SFS_VOLNAME_SIZE 32 /* max length of volume name */
```

Refer to: `kern/include/kern/sfs.h`,

Particular line:

```
char sp_volname[SFS_VOLNAME_SIZE]; /* Name of this volume */
```

Refer to: `userland/sbin/mksfs/mksfs.c`

Particular line: `if (strlen(volname) >= SFS_VOLNAME_SIZE) {`

**Question 8. How many direct blocks does an SFS file have?**

**Answer 8:** 15 direct blocks

Refer to: `kern/include/kern/sfs.h`

Particular line:

```
#define SFS_NDIRECT          15          /* # of direct blocks in inode */
```

**Question 9. What is the standard interface to a file system (i.e., what functions must you implement to implement a new file system)?**

Defined in: `kern/include/fs.h`

切换行号显示

```
1 struct fs_ops {
2     int (*fsop_sync)(struct fs *); //
3     const char *(*fsop_getvolname)(struct fs *); //
4     int (*fsop_getroot)(struct fs *, struct vnode **); //
5     int (*fsop_unmount)(struct fs *); //
6 };
```

To return a useful root vnode in `fs_getroot` you also need to implement all the filesystem specific operations on the vnode, i.e. implement the functions defined in `kern/include/vnode.h`. For SFS, these are defined in `kern/fs/sfs/sfs_vnops.c`

**Question 10. What function puts a thread to sleep?**

`void wchan_sleep(struct wchan *wc, struct spinlock *lk);` as defined in `kern/include/wchan.h`

Specifically, a thread is switched to `S_SLEEP` state by calling

`thread_switch(S_SLEEP, wc, lk);` as in the implementation in `wchan_sleep()` in `kern/thread/thread.c`

==

Isn't this in `thread.c`? Found by tracing through `thread.h` > `struct threadstate_t` > `thread_switch` (`thread.c`) > `wchan_sleep`

**Question 11. How large are OS/161 pids?**

Pid in OS/161 is of type `pid_t` as defined in `kern/include/types.h`, which is then defined as a **32 bit signed integer** in `kern/include/kern/types.h`.

According to `kern/include/limits.h` and `kern/include/kern/limits.h`, a maximum value for process id is set as 32767 and user-level processes can only take a pid that is greater or equal than 2. However, these are not enforced anywhere in current implementation of OS/161.

Particular lines:

切换行号显示

```

1 // in kern/include/kern/types.h:77
2 typedef __i32 __pid_t;      /* Process ID */
3
4 // in kern/include/types.h:127
5 typedef __pid_t pid_t;
6
7 // in kern/include/limits.h:44
8 #define PID_MIN      __PID_MIN
9 #define PID_MAX      __PID_MAX
10
11 // in kern/include/kern/limits.h:76
12 /* Min value for a process ID (that can be assigned to a user process) */
13 #define __PID_MIN      2
14
15 /* Max value for a process ID (change this to match your implementation)
*/
16 #define __PID_MAX      32767
17

```

**Question 12. What operations can you do on a vnode?**

切换行号显示

```

1 struct vnode_ops {
2     unsigned long vop_magic;      /* should always be VOP_MAGIC */
3
4     int (*vop_eachopen)(struct vnode *object, int flags_from_open);
5     int (*vop_reclaim)(struct vnode *vnode);
6
7
8     int (*vop_read)(struct vnode *file, struct uio *uio);
9     int (*vop_readlink)(struct vnode *link, struct uio *uio);
10    int (*vop_getdirent)(struct vnode *dir, struct uio *uio);
11    int (*vop_write)(struct vnode *file, struct uio *uio);
12    int (*vop_ioctl)(struct vnode *object, int op, userptr_t data);
13    int (*vop_stat)(struct vnode *object, struct stat *statbuf);
14    int (*vop_gettype)(struct vnode *object, mode_t *result);
15    bool (*vop_isseekable)(struct vnode *object);
16    int (*vop_fsync)(struct vnode *object);
17    int (*vop_mmap)(struct vnode *file /* add stuff */);
18    int (*vop_truncate)(struct vnode *file, off_t len);
19    int (*vop_namefile)(struct vnode *file, struct uio *uio);
20
21
22    int (*vop_creat)(struct vnode *dir,
23                    const char *name, bool excl, mode_t mode,
24                    struct vnode **result);
25    int (*vop_symlink)(struct vnode *dir,
26                      const char *contents, const char *name);
27    int (*vop_mkdir)(struct vnode *parentdir,
28                    const char *name, mode_t mode);
29    int (*vop_link)(struct vnode *dir,
30                   const char *name, struct vnode *file);
31    int (*vop_remove)(struct vnode *dir,
32                     const char *name);
33    int (*vop_rmdir)(struct vnode *dir,
34                    const char *name);
35
36    int (*vop_rename)(struct vnode *vn1, const char *name1,
37                     struct vnode *vn2, const char *name2);
38
39
40    int (*vop_lookup)(struct vnode *dir,
41                    char *pathname, struct vnode **result);

```

```

42  int (*vop_lookparent)(struct vnode *dir,
43                        char *pathname, struct vnode **result,
44                        char *buf, size_t len);
45 };

```

Refer to: *kern/include/vnode.h*

### Question 13. What is the maximum path length in OS/161?

切换行号显示

```

1  /* Longest full path name */
2  #define __PATH_MAX      1024
3

```

Refer to *kern/include/limits.h:42* and *kern/include/kern/limits.h:63*

Maximum path can be of length 1024 excluding null-terminator as used in *kern/main/menu.c:203*

### Question 14. What is the system call number for a reboot?

I think it is 119

Particular line: `#define SYS_reboot 119`

Refer to *kern/include/kern/syscall.h:197*

### Question 15. Where is `STDIN_FILENO` defined?

It's defined as 0 in *kern/include/kern/unistd.h:34*.

Particular line: `#define STDIN_FILENO 0 /* Standard input */`

Hint: Use Ctrl+Shift+F in vscode can help you find everything you want in the whole code base.

### Question 16. What does `kmain()` do?

Seems to be the kernel main function, ie the first thing that happens when you run the os.

Kernel main. Boot up, then fork the menu thread; wait for a reboot request, and then shut down.

Refer to *kern/main/main.c:212*

### Question 17. What is the difference between `splhigh` and `spl0`?

- `spl0()` sets IPL to 0, enabling all interrupts.
- `splhigh()` sets IPL to the highest value, disabling all interrupts.

IPL = Interrupt Priority level I believe

Refer to *kern/include/spl.h:94* and *kern/thread/spl.c:128*

### Question 18. What does `splx` return?

An int, representing the old interrupt state/spl level for the current thread.

Refer to *kern/thread/spl.c:128*

### Question 19. What is a zombie?

Zombies are threads that have exited but still need to have `thread_destroy` called on them.

Refer to *kern/include/thread.h:64*

### Question 20. What does a device name in OS/161 look like?

```

/*
 * The name of a device is always just "device:". The VFS
 * layer puts in the device name for us, so we don't need to
 * do anything further.
 */

```

The bit after device: is

```

* kd_name      - Name of device (eg, "lhd0"). Should always be set to
*                a valid string.

```

Refer to *kern/vfs/device.c:281* and *kern/vfs/vfslst.c:50*

- `vfs_adddev` - Add a device to the VFS named device list. If
  - `MOUNTABLE` is zero, the device will be accessible as `"DEVNAME:"`. If the mountable flag is set, the device will be accessible as `"DEVNAMEraw:"` and mountable under the name `"DEVNAME"`. Thus, the console, added with `MOUNTABLE` not set, would be accessed by pathname as `"con:"`, and `lhd0`, added with mountable set, would be accessed by pathname as `"lhd0raw:"` and mounted by passing `"lhd0"` to `vfs_mount`.

**Question 21. What does a raw device name in OS/161 look like?**

```

*
* kd_rawname - Name of raw device (eg, "lhd0raw"). Is non-NULL if and
*                only if this device can have a filesystem mounted on
*                it.

```

Refer to *kern/vfs/vfslst.c:53*

**Question 22. What lock protects the vnode reference count?**

A spinlock that is a member of `struct vnode` as defined in *kern/include/vnode.h:51*.

Particular lines:

切换行号显示

```

1 struct vnode {
2     int vn_refcount;           /* Reference count */
3     struct spinlock vn_countlock; /* Lock for vn_refcount */
4     // ...
5 };
6
7 //=====
8
9 // in kern/vfs/vnode.c:84
10 spinlock_acquire(&vn->vn_countlock);
11 vn->vn_refcount++;
12 spinlock_release(&vn->vn_countlock);

```

**Question 23. What device types are currently supported?**

I think its:...

A device is either a "block device" or a "character device".

Since the documented return value of

```
static int dev_gettype(struct vnode *v, mode_t *ret) is:
```

- Return the type. A device is a "block device" if it has a known length. A device that generates data in a stream is a "character device".

Refer to *kern/vfs/device.c:195*

